

---

---

컴퓨터 공학 개론  
**Lecture 6**

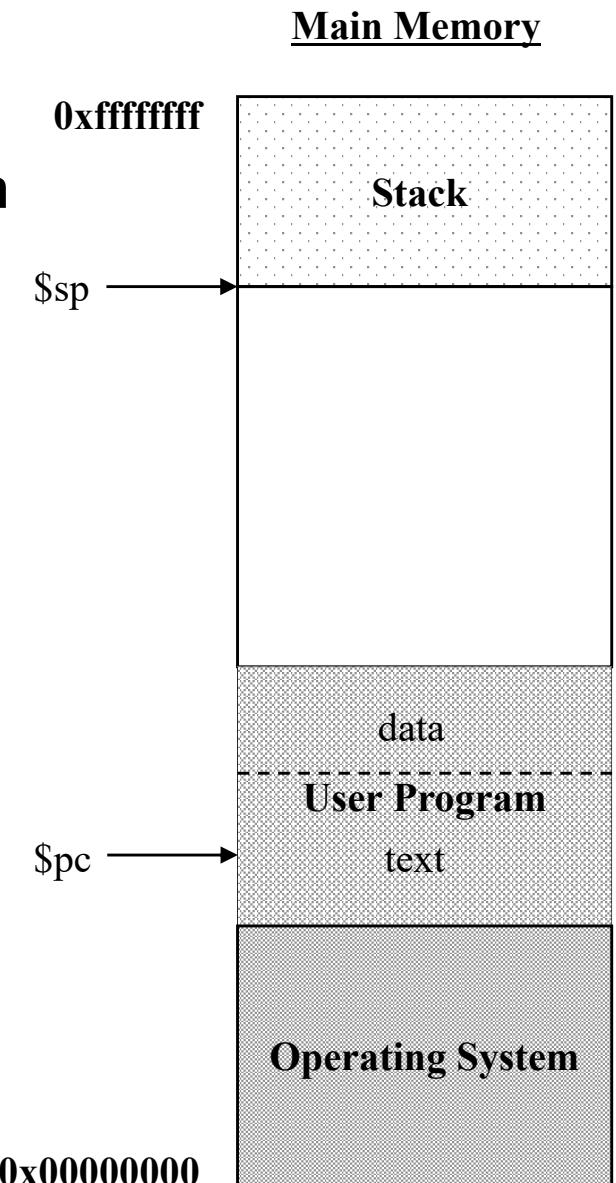
**Instructions: Language of the Computer**

**2017**

**김태성**

# Memory Allocation

- \$pc
  - Program Counter
  - register for the address of current instruction
- Stack
  - a last-in-first-out(LIFO) data structure
  - memory reserved for spilling registers
  - grows downward
  - *push: place data onto the stack*
  - *pop: remove data from the stack*
- \$sp
  - Stack Pointer
  - register for the stack
  - most recently allocated address in the stack

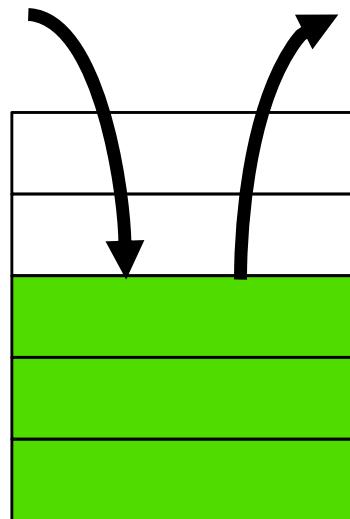


# Stack vs. FIFO

---

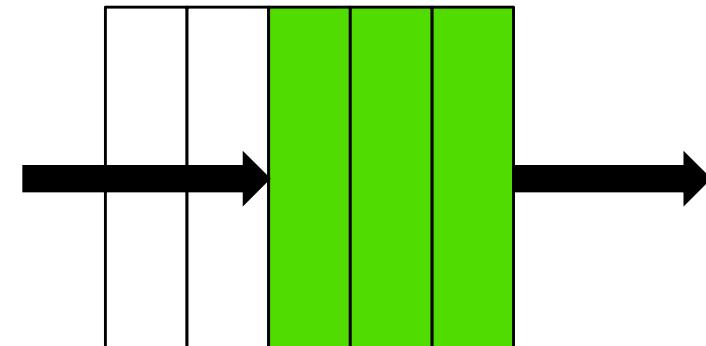
- Stack

- Last-In-First-Out(LIFO) data structure



- FIFO

- First-In-First-Out data structure



# Procedure Calls

---

- In the execution of a procedure, the program must follow these six steps:

1. Place parameters in a place where the procedure can access them
2. Transfer control to the procedure
3. Acquire the storage needed for the procedure
4. Perform the procedure operations
5. Place the result value in a place where the calling program can access it
6. Return control to the caller

} Caller side

} Callee side

```
main(...){  
...  
...  
flag = bar(a, b);  
....  
}  
...  
  
int bar(x,y){  
...  
...  
return(result);  
}
```

# Procedure Calls

---

1. Place parameters in a place where the procedure can access them
  - \$a0 - \$a3: argument registers to pass parameters
2. Transfer control to the procedure
  - \$ra: register for return address
  - jal ProcedureAddress        # saves the address of the next instruction in \$ra  
                                    # and jump to “ProcedureAddress”
3. Acquire the storage needed for the procedure
  - save registers into stack
4. Perform the procedure operation
5. Place the result value in a place where the calling program can access it
  - \$v0 - \$v1: two registers for return value
6. Return control to the caller
  - restore saved registers
  - jr \$ra                      # ‘jump register’: jump to the address stored in \$ra

# Calling a Procedure

---

- C Program:

```
int bar (int g, int h, int i, int j){  
    int f;  
    f = ( g + h ) - ( i + j );  
    return f;  
}
```

- MIPS Assembly Program:

	# step 1, 2 are executed on the caller side
	# compiler assigned \$a0, \$a1, \$a2, \$a3, and \$s0
	# for g, h, i, j, and f respectively
bar:	# procedure starts here executing step 3
addi \$sp, \$sp, -12	# adjust stack to make room for 3 items
sw \$t1, 8(\$sp)	# save register \$t1 for use afterwards
sw \$t0, 4(\$sp)	# save register \$t0 for use afterwards
sw \$s0, 0(\$sp)	# save register \$s0 for use afterwards

# Calling a Procedure (cont.)

---

# step 4: main body of the procedure

```
add $t0, $a0, $a1 # register $t0 contains ( g + h )
add $t1, $a2, $a3 # register $t1 contains ( i + j )
sub $s0, $t0, $t1 # f gets ( $t0 - $t1 ), which is ( g + h ) - ( i + j )
```

add \$v0, \$s0, \$zero      # step 5: returns f ( \$v0 = \$s0 + 0 )

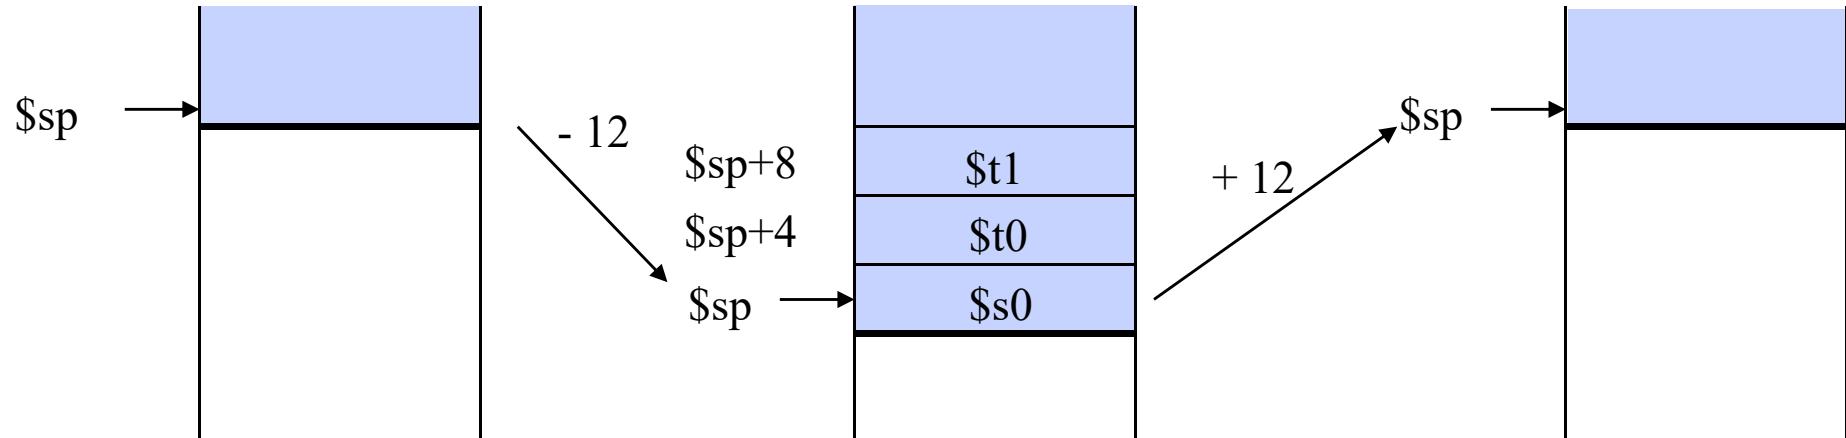
# step 6:

```
lw $s0, 0($sp)      # restore register $s0 for caller
lw $t0, 4($sp)      # restore register $t0 for caller
lw $t1, 8($sp)      # restore register $t1 for caller
addi $sp, $sp, 12    # adjust stack to delete 3 items
jr $ra              # jump back to calling routine
```

# Stack Movement

---

*High Address*



*Low Address*

**Before the procedure call**

**During the procedure call**

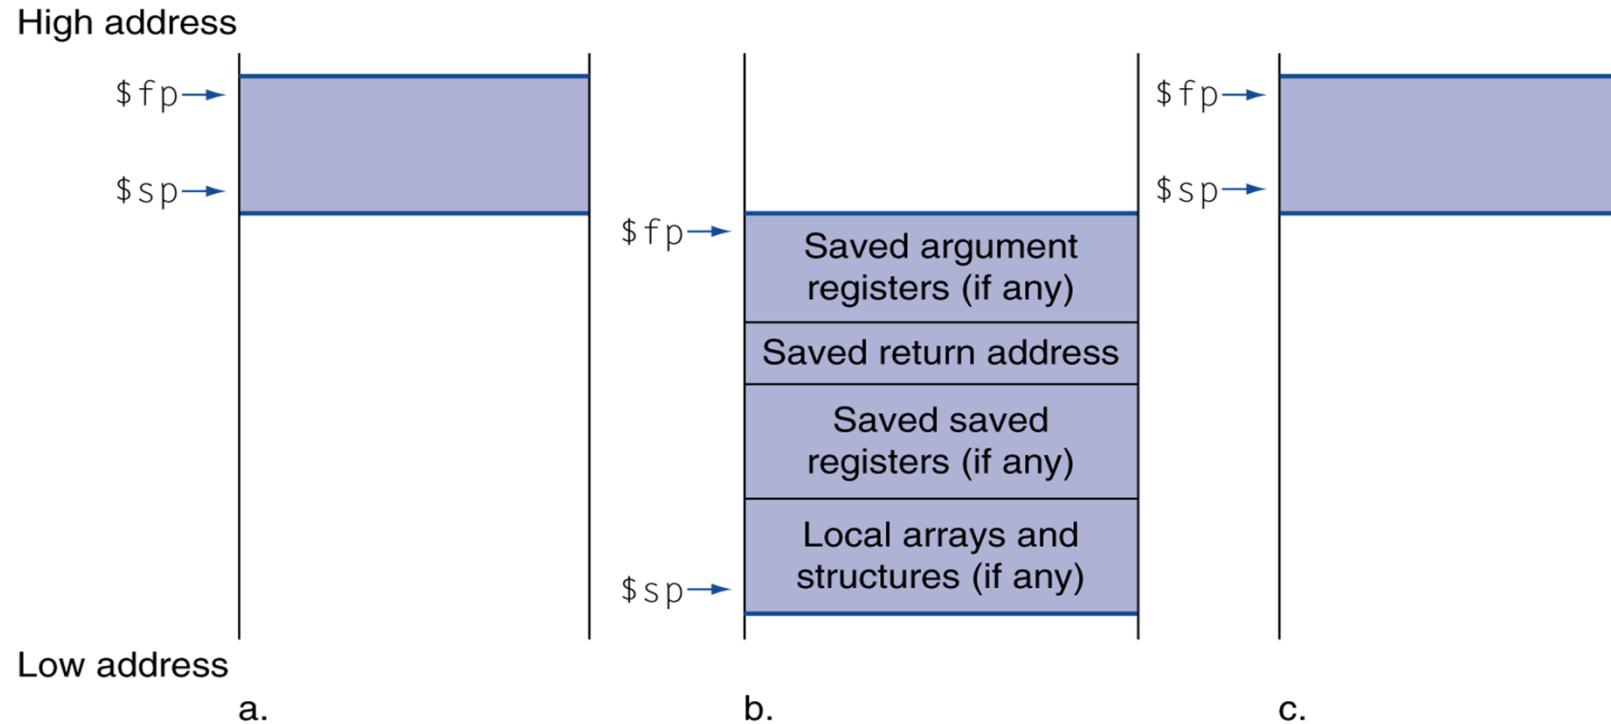
**After the procedure call**

# Nested Procedures

---

- **Procedures that call other procedures**
- **For nested call, caller needs to save on the stack:**
  - Its return address (\$ra)
  - Any arguments and temporaries needed after the call (\$a0~\$a3)
- **Restore from the stack after the call**

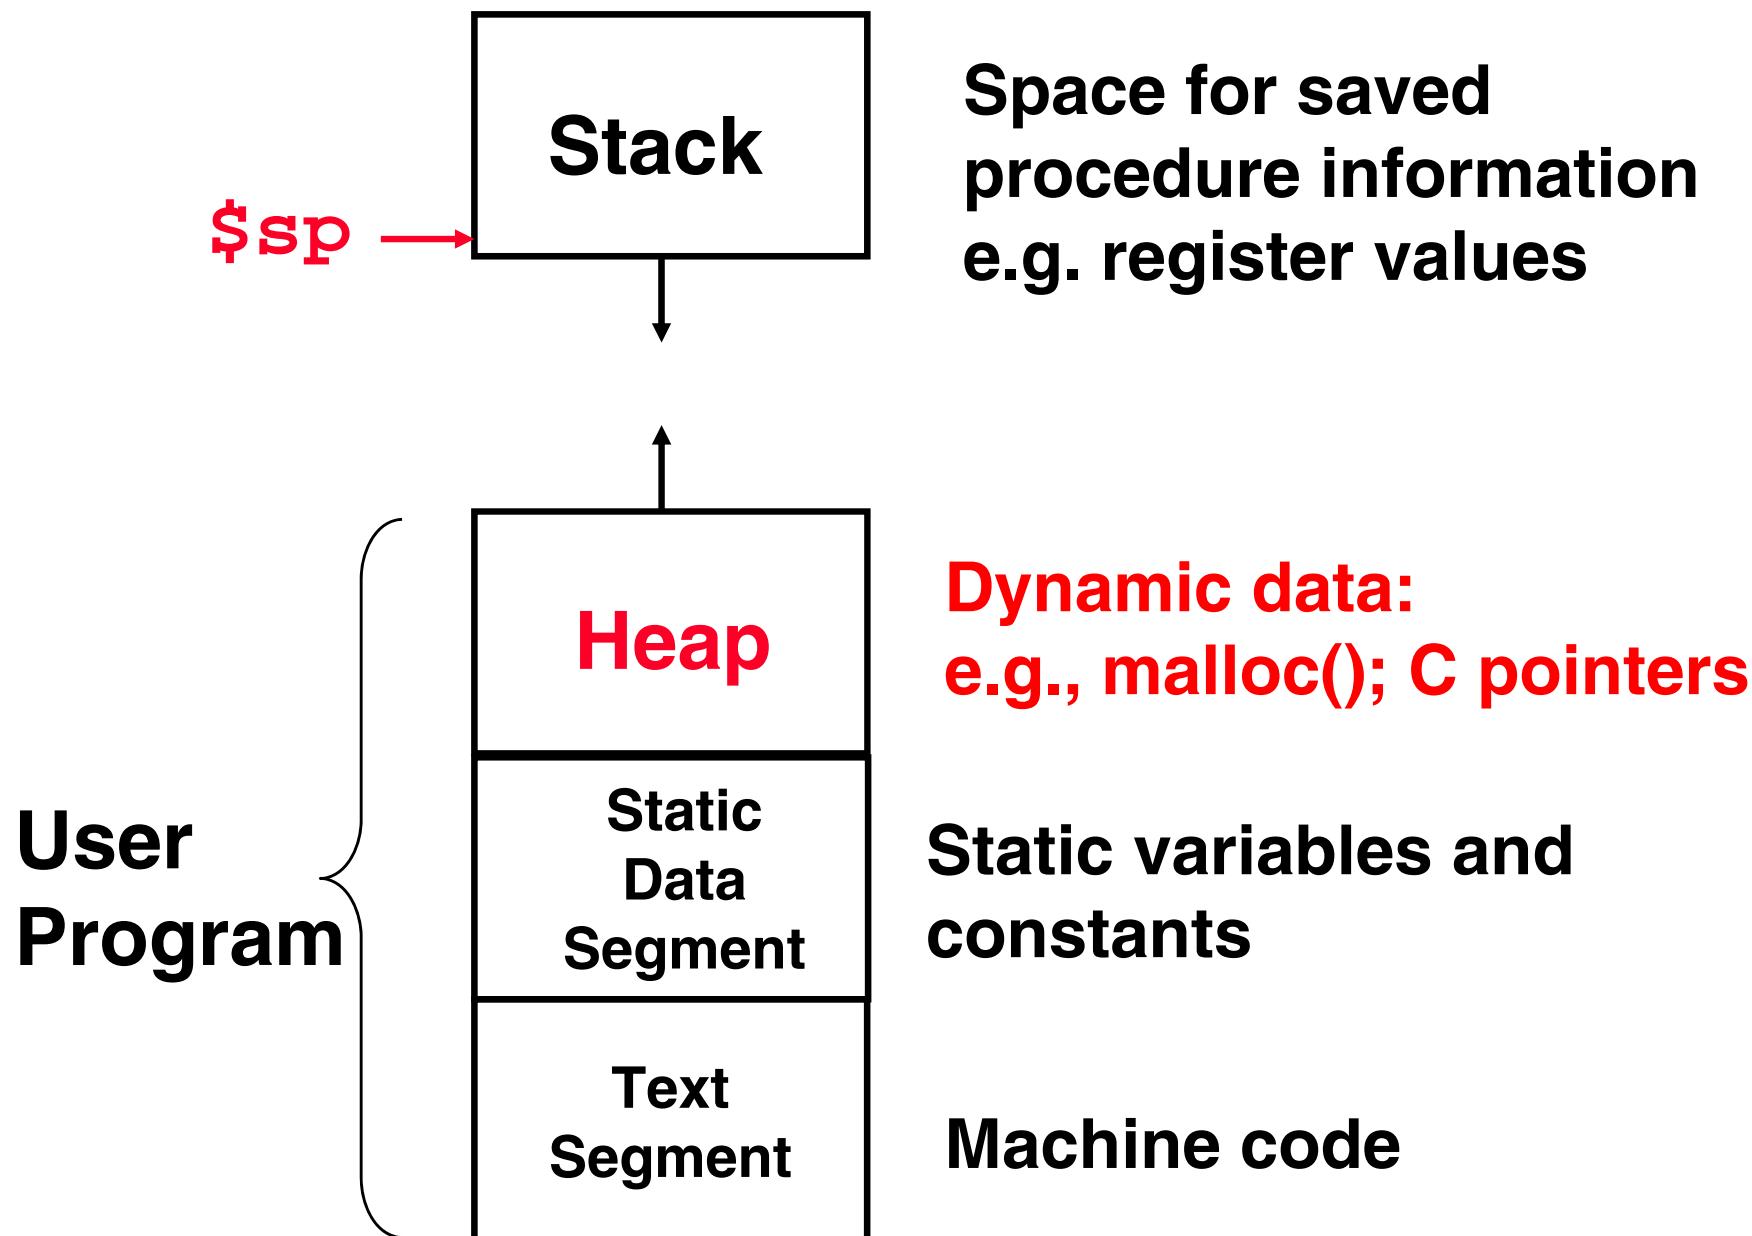
# Local Data on the Stack



- ° Local data allocated by callee
  - e.g., C automatic variables (<-> external variables)
- ° Procedure frame (activation record)
  - Used by some compilers to manage stack storage
  - **\$fp: frame pointer**

# Inside of User Program

---



# MIPS Register Summary

---

Name	Register Number	Usage
\$zero	0	The constant value 0
\$at	1	<i>Reserved for assembler</i>
\$v0 - \$v1	2-3	Return value
\$a0 - \$a3	4-7	Arguments
\$t0 - \$t7	8-15	Temporaries
\$s0 - \$s7	16-23	Static variables
\$t8 - \$t9	23-25	More temporaries
\$k0 - \$k1	26-27	<i>Reserved for OS kernel</i>
\$gp	28	Global pointer
\$sp	29	Stack pointer
\$fp	30	Frame pointer
\$ra	31	Return address

# Kilo, Mega, Giga, Tera, Peta, Exa, Zetta, Yotta

---

Name	Abbr	Factor	SI (International System of Units) size
Kilo	K	$2^{10} = 1,024$	$10^3 = 1,000$
Mega	M	$2^{20} = 1,048,576$	$10^6 = 1,000,000$
Giga	G	$2^{30} = 1,073,741,824$	$10^9 = 1,000,000,000$
Tera	T	$2^{40} = 1,099,511,627,776$	$10^{12} = 1,000,000,000,000$
Peta	P	$2^{50} = 1,125,899,906,842,624$	$10^{15} = 1,000,000,000,000,000$
Exa	E	$2^{60} = 1,152,921,504,606,846,976$	$10^{18} = 1,000,000,000,000,000,000$
Zetta	Z	$2^{70} = 1,180,591,620,717,411,303,424$	$10^{21} = 1,000,000,000,000,000,000,000$
Yotta	Y	$2^{80} = 1,208,925,819,614,629,174,706,176$	$10^{24} = 1,000,000,000,000,000,000,000,000$

# Kilo, Mega, Giga, Tera, Peta, Exa, Zetta, Yotta

---

- Common use prefixes in computer engineering
- Confusing! Common usage of “kilobyte” means 1024 bytes, but the “correct” SI value is 1000 bytes
- $4\text{Gb DRAM} = 4 \times 2^{30} \text{ bits} = 4,294,967,296 \text{ bits}$
- $1.3\text{GB AVI file} = 1.3 \times 2^{30} \text{ Bytes} = 1,395,864,371 \text{ Bytes}$
- Hard Disk manufacturers & Telecommunications use SI factors, so what is advertised as a 300 GB drive will actually only hold about  $280 \times 2^{30} \text{ bytes}$ , and a 1 Mbit/s connection transfers  $10^6 \text{ bps}$ .

# Handling Characters

ASCII: American Standard Code for Information Interchange

Character	ASCII Value
...	...
?	63 (011 1111)
@	64 (100 0000)
A	65 (100 0001)
B	66 (100 0010)
...	...

- **How to Represent Characters?**

- English alphabet: 26 letters
- upper/lower case + symbols + formatting ( < 128 char.)
  - => can be represented in 7 bits (“ASCII”)
  - => use 1 byte(8 bits) to store one character

- **MIPS instructions for loading/storing character data**

- **lb:** “load byte” loads a byte from memory and place it in the rightmost 8 bits of a register
  - e.g. `lb $t0, 0($s1) // read byte from memory`
- **sb:** “store byte” takes a byte from the rightmost 8 bits of a register and writes it to memory
  - e.g. `sb $t0, 0($s1) // write byte to memory`

# Handling Characters

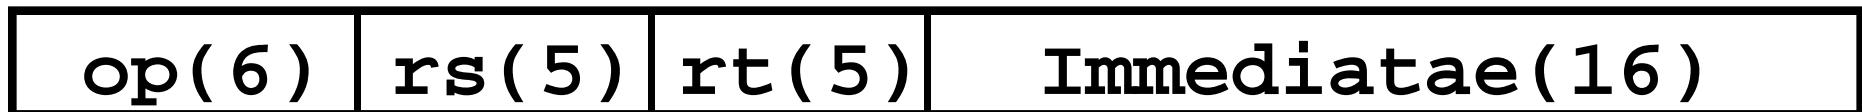
---

- Standard code to cover all the world’s languages
  - => 8,16,32 bits (“Unicode”, [www.unicode.com](http://www.unicode.com))
  - 16-bit encoding (UTF-16) is default
    - UTF: Unicode Transformation Formats
  - UTF-8 (variable-length encoding) keeps the ASCII subset as 8 bits and use 16~32 bits for the other characters
  - UTF-32 uses 32 bits per characters
- MIPS instructions for loading/storing 16-bit character data
  - lh: “load half” loads a halfword from memory and place it in the rightmost 16 bits of a register
    - e.g. lh \$t0, 0(\$s1) // read halfword (16 bits) from memory
  - sh: “store half” takes a halfword from the rightmost 16 bits of a register and writes it to memory
    - e.g. sh \$t0, 0(\$s1) // write halfword (16 bits) to memory

# Constants

---

- Small constants are used quite frequently (50% of operands)  
e.g.,     $A = A + 5;$   
               $B = B + 1;$   
               $C = C - 18;$
- MIPS Instructions:  
`addi $29, $29, 4`  
`slti $8, $18, 10`  
`andi $29, $29, 6`  
`ori $29, $29, 4`
- These are I-type instructions:

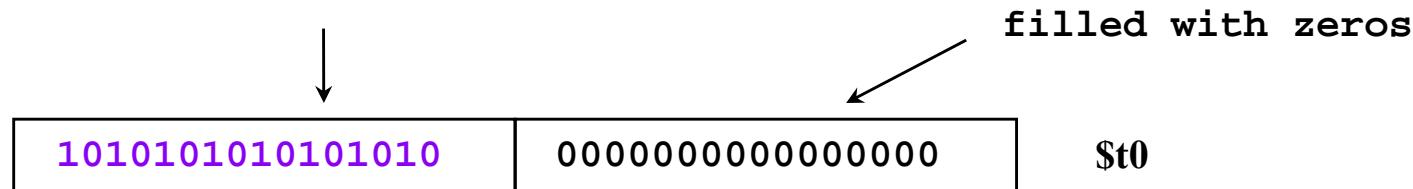


# How about larger constants?

---

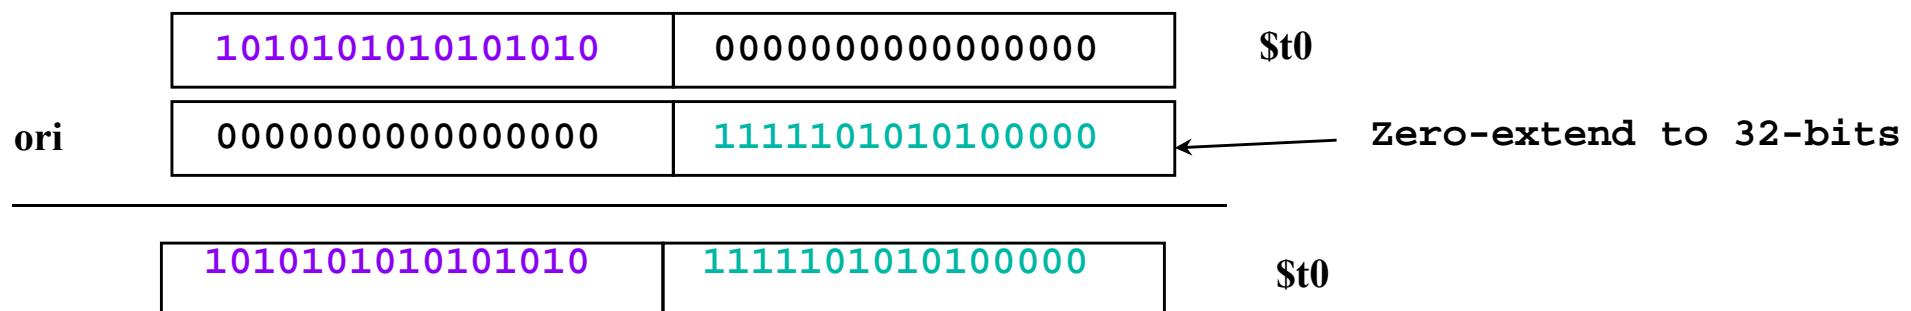
- We load a 32 bit constant into a register first, then operate on the register
- Must use two instructions, new "load upper immediate" instruction

**lui \$t0, 1010101010101010**



- Then must get the lower order bits right, i.e.,

**ori \$t0, \$t0, 1111101010100000**



## Load Upper Immediate (lui)

---

- “lui” is an I-type Instruction
- lui \$t0, 10101010101010
  - interpreted as lui \$t0, \$zero, 10101010101010
  - similar to addi \$t0, \$zero, 10101010101010  
except the immediate operand is considered as the upper 16 bits of a constant

# Addresses in Branches and Jumps

---

- **Instructions:**
  - **bne \$t4,\$t5,Label**      Next instruction is at Label if  $\$t4 \neq \$t5$
  - **beq \$t4,\$t5,Label**      Next instruction is at Label if  $\$t4 = \$t5$
  - **j Label**                  Next instruction is at Label
- **Formats:**

I	op(6)	rs(5)	rt(5)	16 bit address
J	op(6)	26 bit address		

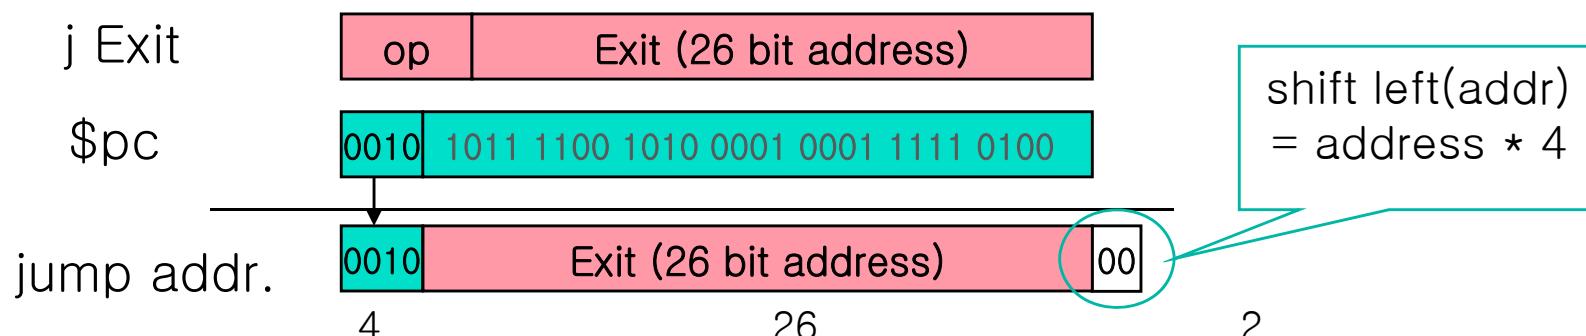
  - Addresses are not 32 bits (either 16 bits or 26 bits)
    - How do we handle this?

# Addresses in Branches and Jumps

- **Branches: PC-relative Addressing**
  - Could specify a register (as in `lw` and `sw`) and add it to address
  - Use program counter(`$pc`) as a base address
  - e.g. `beq $s1, $s2, 25`      `# if ($s1 == $s2) go to ($pc+4)+(4*25)`
  - `# else go to ($pc+4) (next instruction)`

cf. `lw $t1, 16($s0)`

- **Jumps: Pseudodirect Addressing**
  - Jump instructions just use high order bits of `$pc`
  - the jump address is the 26 bits of the instruction concatenated with the upper 4 bits of the `$pc`



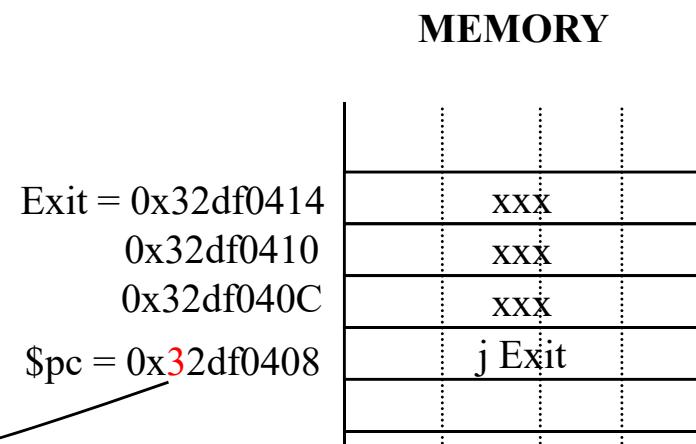
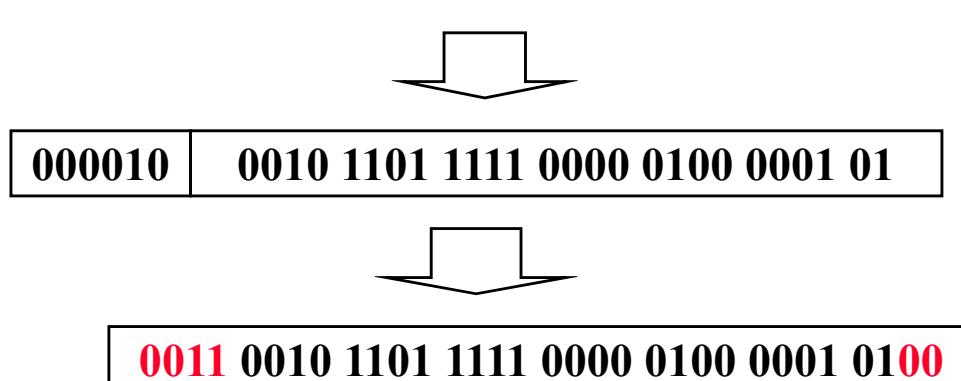
# More about Jump

- Current address(\$pc) = 0x32df0408  
(0011 0010 1101 1111 0000 0100 0000 1000)
- Destination address(Exit) = 0x32df0414  
(0011 0010 1101 1111 0000 0100 0001 0100)
- Instruction ‘j Exit’

000010	0011	0010 1101 1111 0000 0100 0001 0100
--------	------	------------------------------------

6 bits + 32 bits = 38 bits => Too many!

- Last 2 bits are always zero because each instruction consists of 4 bytes.
- First 4 bits remains the same with those in \$pc as long as the program is placed within 256MB boundary



# Target Addressing Example

- Loop code from earlier example
  - Assume Loop at location 80000

Loop: sll \$t1, \$s3, 2  
add \$t1, \$t1, \$s6  
lw \$t0, 0(\$t1)  
bne \$t0, \$s5, Exit  
addi \$s3, \$s3, 1  
j Loop  
Exit: ...

80000	0	0	19	9	4	0
80004	0	9	22	9	0	32
80008	35	9	8			0
80012	5	8	21			2
80016	8	19	19			1
80020	2					20000
80024						

# MIPS Addressing Mode

---

## 1. Register addressing

- operand is a register
- e.g.            add \$s0, \$s1, \$s2 #  $\$s0 = \$s1 + \$s2$

## 2. Base or Displacement addressing

- operand = Mem[register value + constant]
- e.g.            lw \$s1, 100(\$s2)                            #  $\$s1 = \text{Mem}[\$s2 + 100]$

## 3. Immediate addressing

- operand is a constant within the instruction
- e.g.            addi \$sp, \$sp, 12                        #  $\$sp = \$sp + 12$
- slti \$s1, \$s2, 100                        #  $\$s1 = 1 \text{ if } (\$s2 < 100)$
- # else  $\$s1 = 0$

## 4. PC-relative addressing

- address is the sum of PC and a constant in the instruction
- e.g.            beq \$s1, \$s2, 25                        # if ( $\$s1 == \$s2$ ) go to PC+4+(4\*25)

## 5. Pseudodirect addressing

- the jump address is the 26 bits of the instruction concatenated with the upper bits of the \$pc

# Assembler Pseudoinstructions

---

- **Most assembler instructions represent machine instructions one-to-one**
- **Pseudoinstructions: figments of the assembler's imagination**

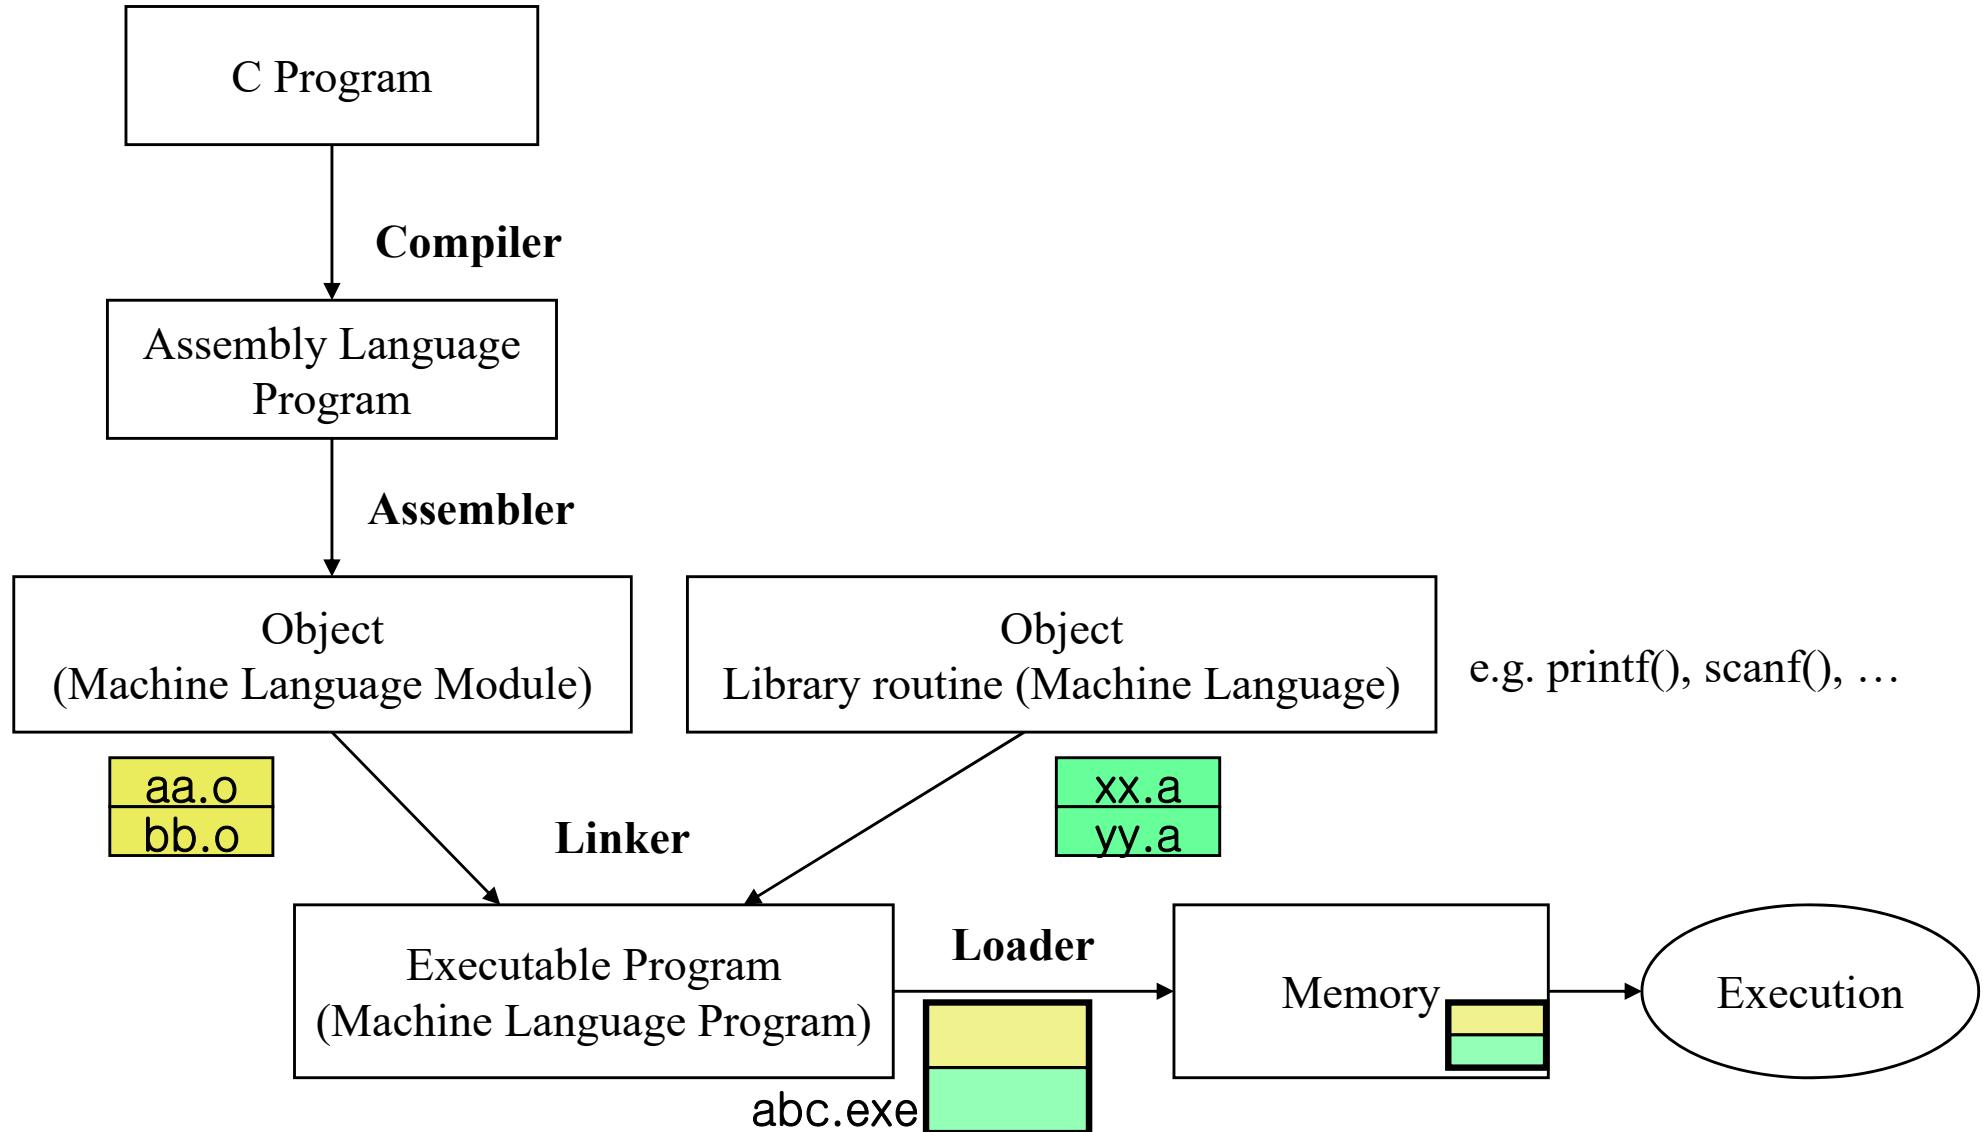
move \$t0, \$t1 → add \$t0, \$zero, \$t1

blt \$t0, \$t1, L → slt \$at, \$t0, \$t1  
bne \$at, \$zero, L

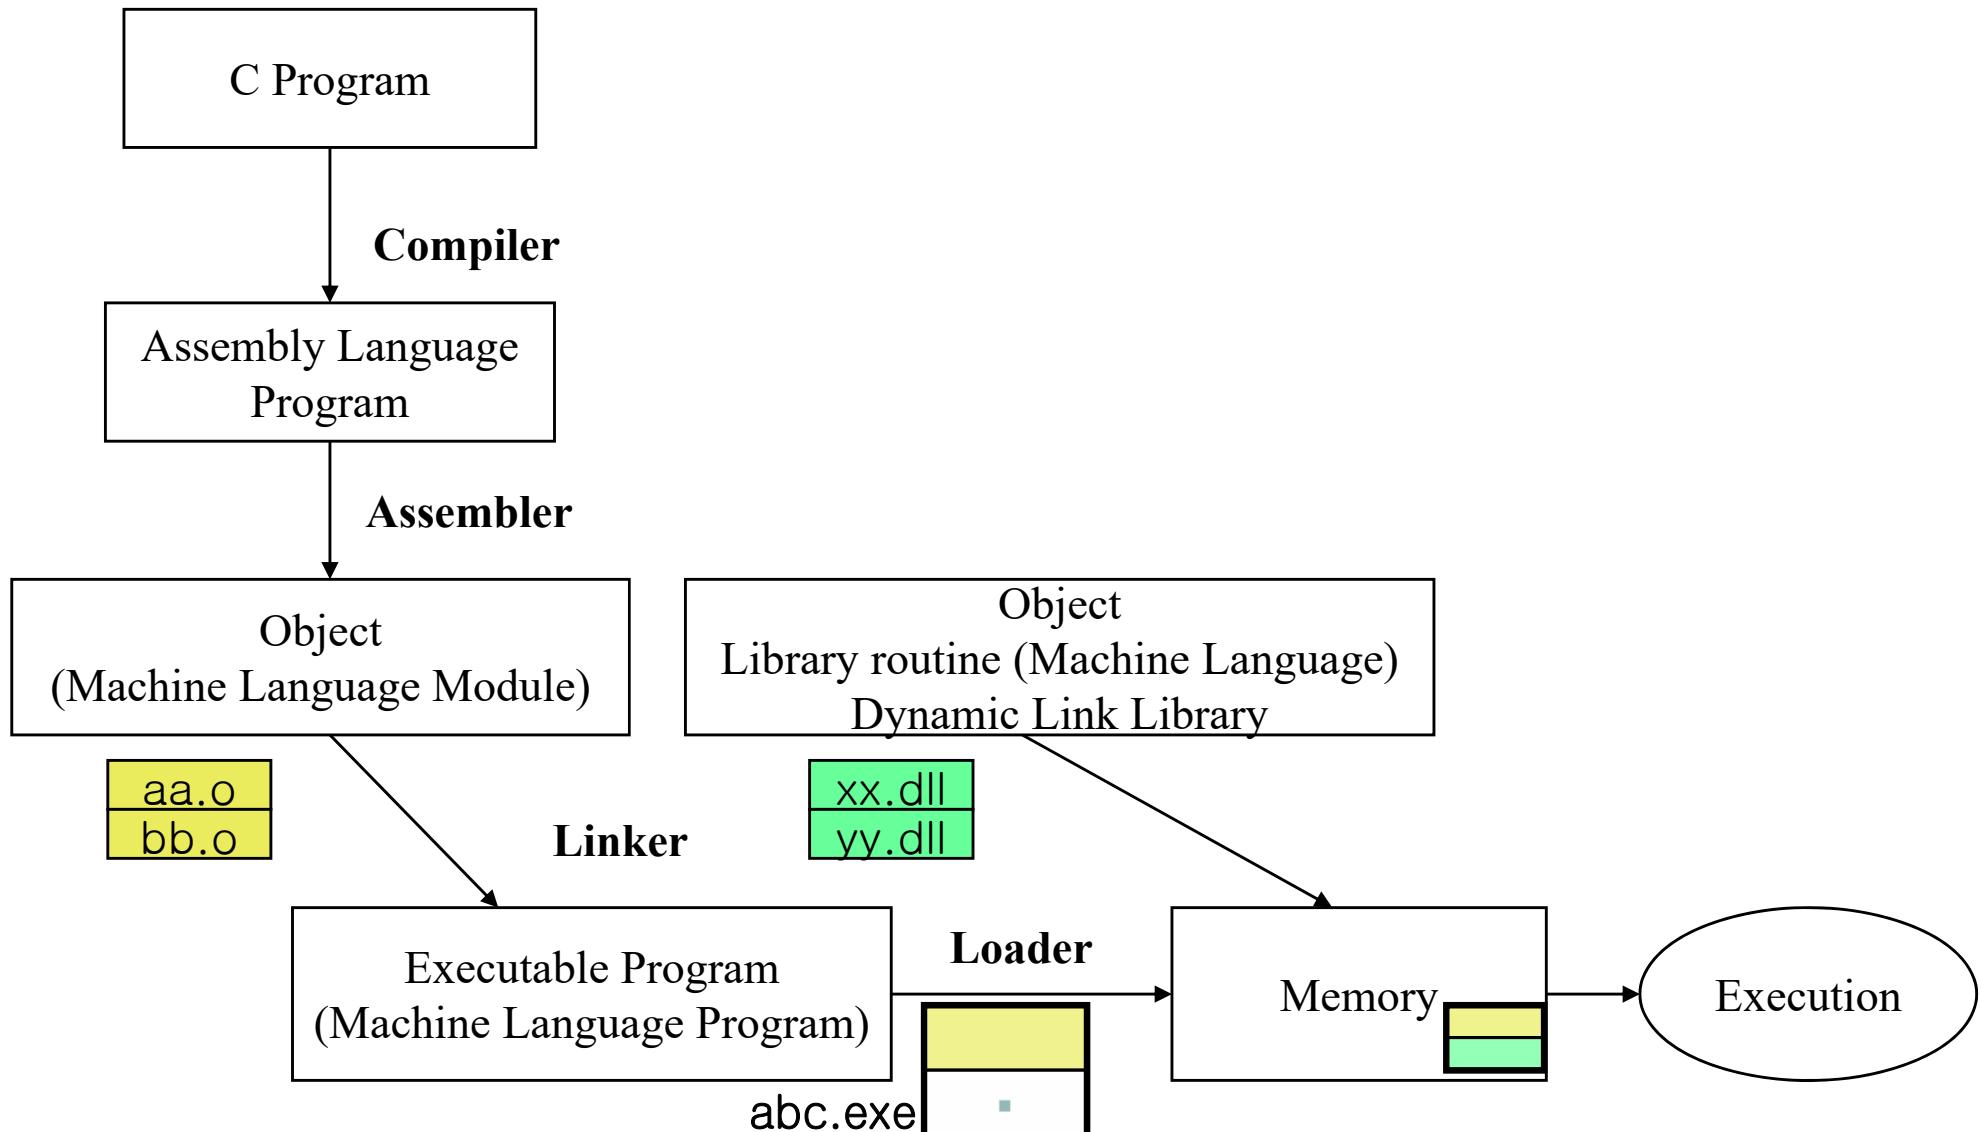
- **\$at (register 1): assembler temporary**

\***figment of imagination:** 상상의 산물

# Starting a Program



# Starting a Program with DLL



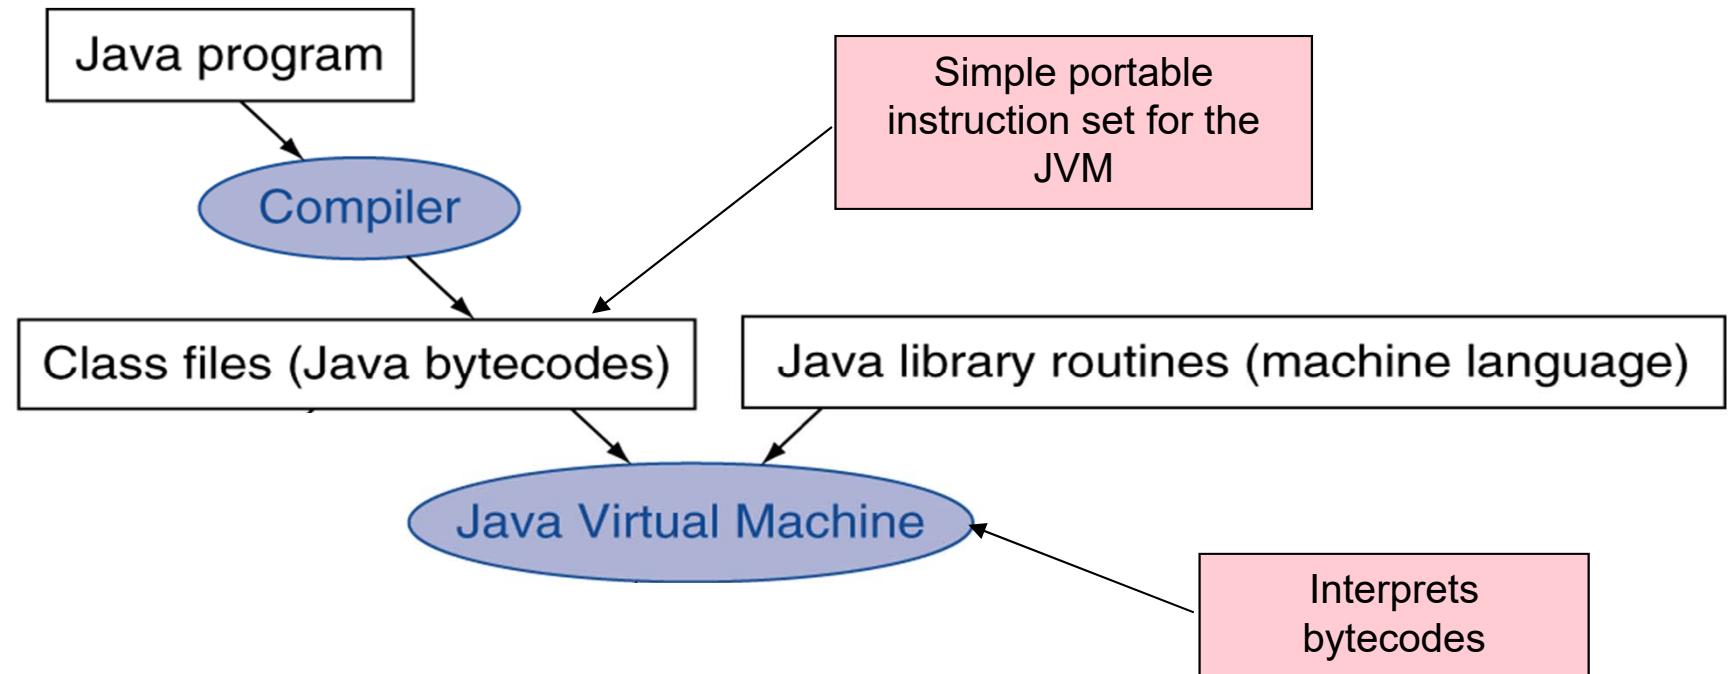
# Dynamically Linked Libraries (DLL)

---

- **library routines are not linked into the executable program**
- **they are linked and loaded when the program is run**
- **library routines can be bug-fixed and updated independently from application programs**
- **can save memory while program is running**
  - **More than one program can share the same dll's**
- **MS Windows makes extensive use of DLLs**

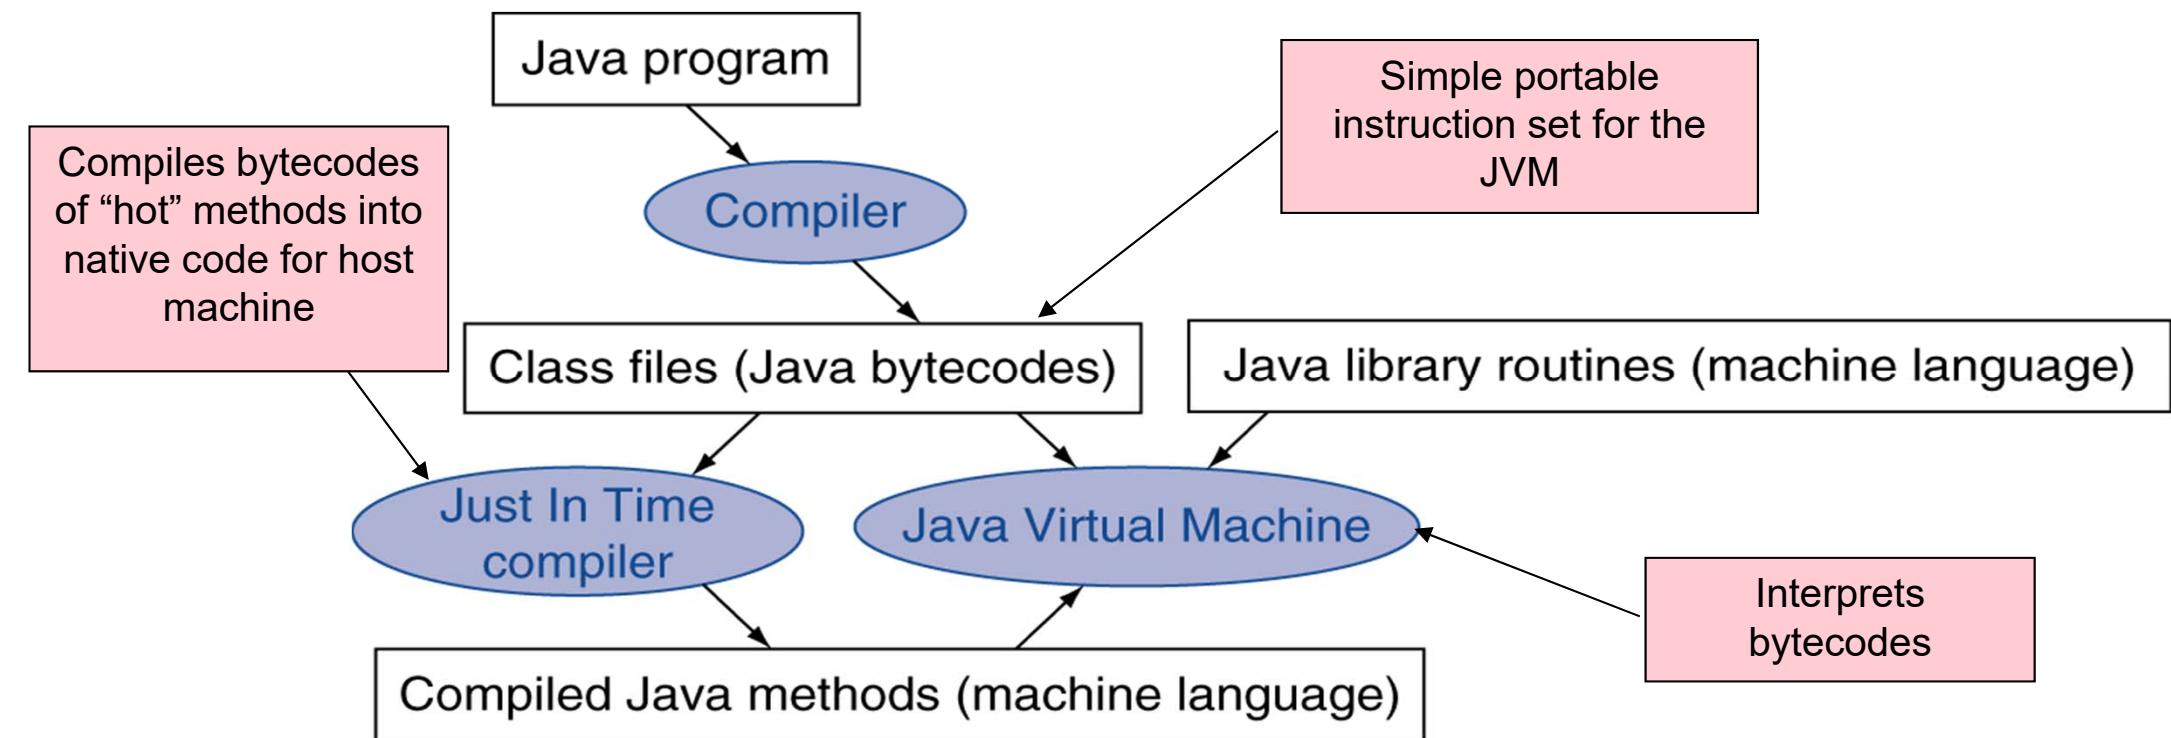
# Starting Java Applications

---



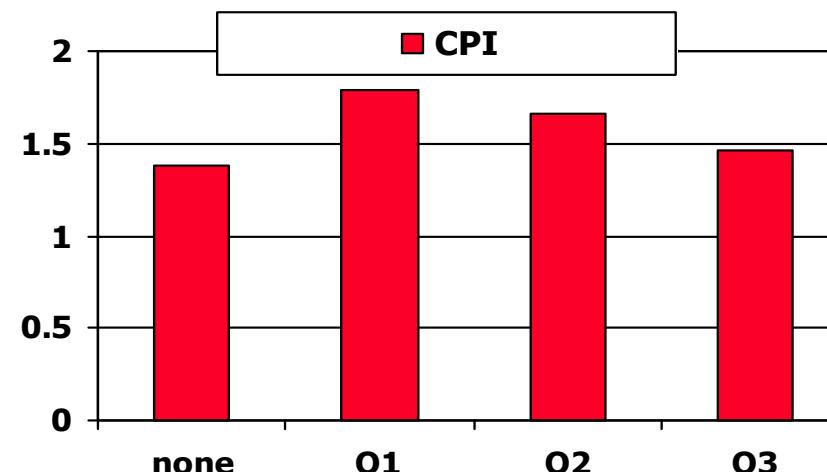
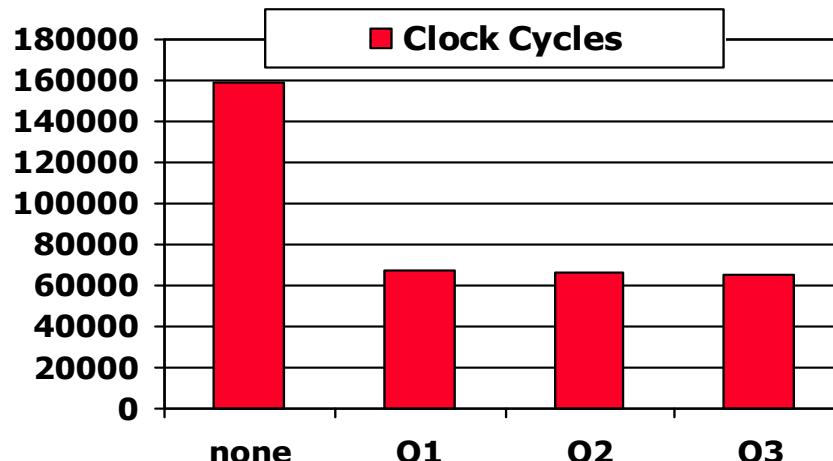
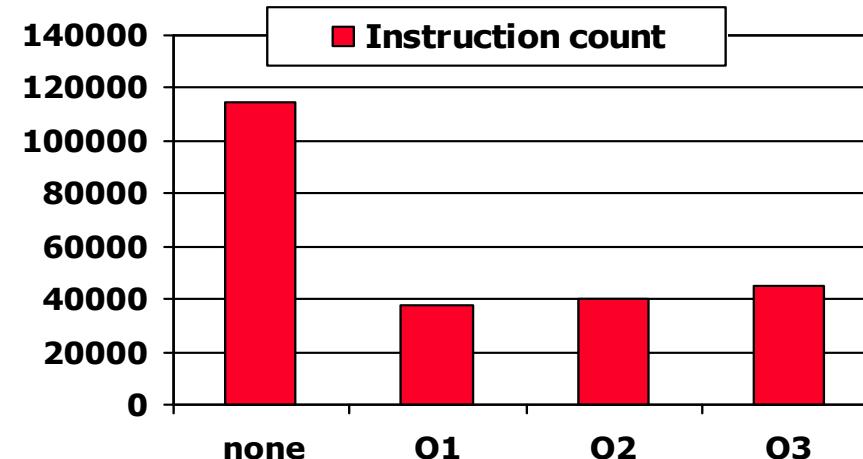
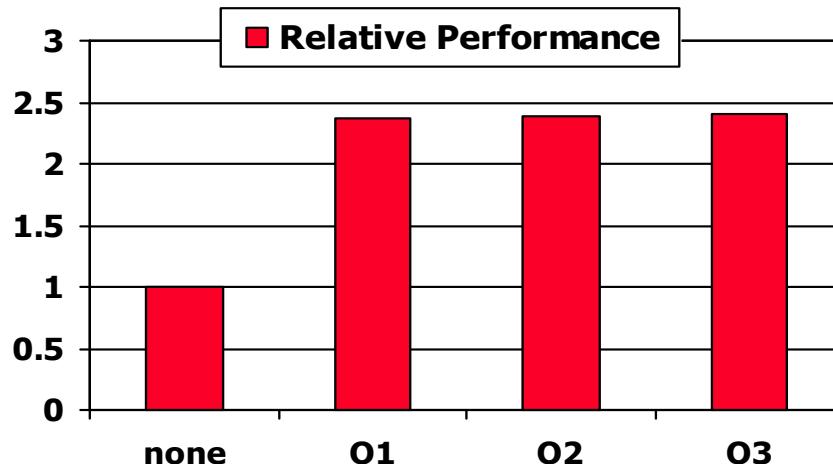
# Starting Java Applications

---



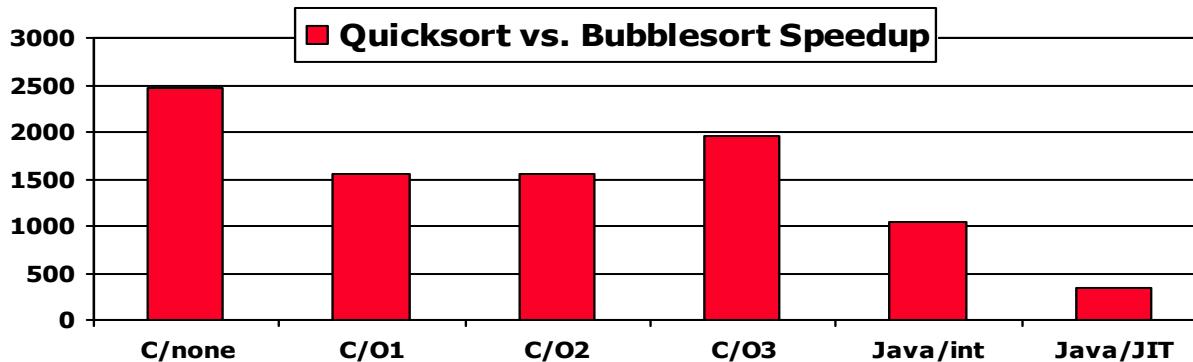
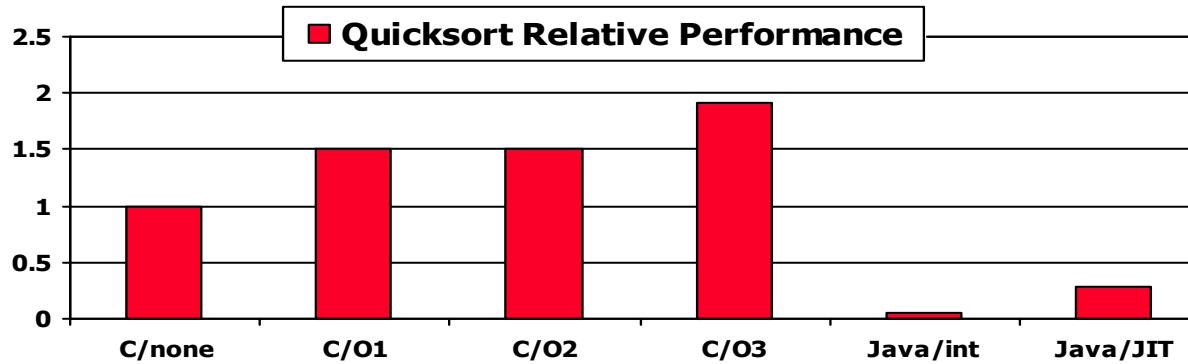
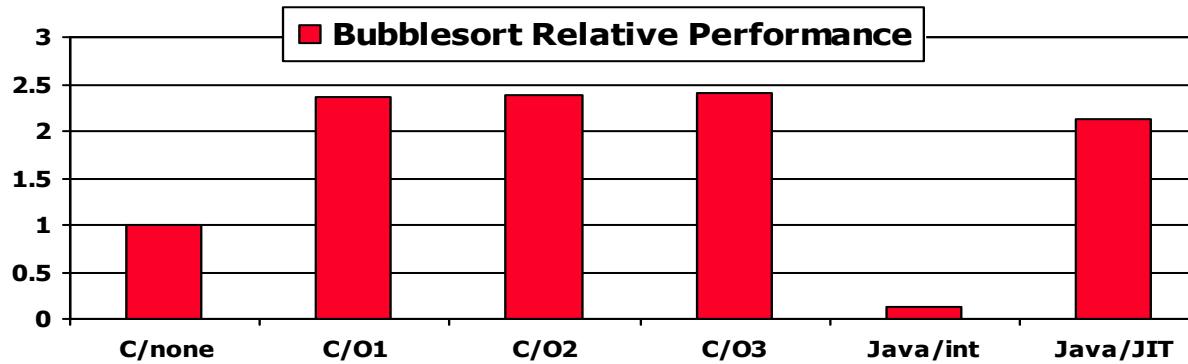
# Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux



# Effect of Language and Algorithm

---



# Lessons Learnt

---

- Instruction count and CPI are not good performance indicators in isolation
- Compiler optimizations are sensitive to the algorithm
- Java/JIT compiled code is significantly faster than JVM interpreted
  - Comparable to optimized C in some cases
- Nothing can fix a dumb algorithm!