
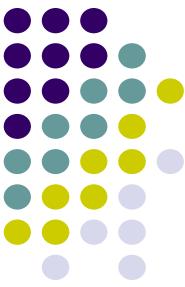


컴퓨터 공학 개론

Lecture 8


2017





Floating Point

- Representation for non-integral numbers
 - Including very small and very large numbers
- Like scientific notation
 - -2.34×10^{56} ← normalized
 - $+0.002 \times 10^{-4}$ ← not normalized
 - $+987.02 \times 10^9$ ← not normalized
- In binary
 - $\pm 1.xxxxxxx_2 \times 2^{yyy}$
- Types **float** and **double** in C
 - vs. **int** and **long**





Recall Scientific Notation

- Number represented as
 - fraction
 - exponent (using radix 10)

- Arithmetic
 - Multiplication, Division
 - multiply/divide fraction
 - add/subtract exponent
 - normalize
 - example: $(5.6 \times 10^{11}) \times (6.7 \times 10^{12}) = (5.6 \times 6.7) \times 10^{(11+12)}$
 $= 37.52 \times 10^{23} = 3.752 \times 10^{24}$


- Addition, Subtraction
 - convert operands to have the same exponent value
 - add/subtract fraction
 - normalize (if needed)
 - example: $(2.1 \times 10^3) + (4.3 \times 10^4) = (0.21 \times 10^4) + (4.3 \times 10^4) = 4.51 \times 10^4$





Floating Point Standard

- Defined by IEEE Std 754–1985
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Two representations
 - Single precision (32-bit)
 - Double precision (64-bit)



IEEE Floating-Point Format

single: 8 bits

double: 11 bits


single: 23 bits

double: 52 bits



$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit ($0 \Rightarrow$ non-negative, $1 \Rightarrow$ negative)
- Normalize significand ($\equiv (1 + \text{Fraction})$)
 - $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1203




Floating-Point Example

- What number is represented by the single-precision float


11000000101000…00

- S = 1
- Exponent = $10000001_2 = 129$
- Fraction = $0.01000\cdots00_2 (= 0 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3} + \dots)$
- $x = (-1)^1 \times (1 + .01_2) \times 2^{(129 - 127)}$
$$= (-1) \times 1.25 \times 2^2$$
$$= -5.0$$



Floating-Point Example

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000\cdots00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 011111111110_2$
- Single: $101111101000\cdots00$
- Double: $1011111111101000\cdots00$




Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
 \Rightarrow actual exponent = $1 - 127 = -126$
 - Fraction: 000 \cdots 00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 \Rightarrow actual exponent = $254 - 127 = +127$
 - Fraction: 111 \cdots 11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$




Double-Precision Range

- Exponents 0000…00 and 1111…11 reserved
- Smallest value
 - Exponent: 00000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000…00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 11111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111…11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$



Floating-Point Precision


- Relative precision
 - all fraction bits are significant
 - Single: approx 2^{-23}
 - Equivalent to $23 \times \log_{10}2 \approx 23 \times 0.3 \approx 6$ decimal digits of precision
 - Double: approx 2^{-52}
 - Equivalent to $52 \times \log_{10}2 \approx 52 \times 0.3 \approx 16$ decimal digits of precision




Infinities and NaNs

- Exponent = 111...1, Fraction = 000...0
 - $\pm\infty$
 - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction \neq 000...0
 - Not-a-Number (NaN)
 - Indicates illegal or undefined result
 - e.g., $0.0 / 0.0$, $\sqrt{-4}$
 - Can be used in subsequent calculations

Range of Numbers Represented by IEEE 754



sign	Exponent	Fraction		
0/1	1111 1111	1111 1111 ... 1111 1111	NaN	
0/1	1111 1111	...		
0/1	1111 1111	0000 0000 ... 0000 0001		
0/1	1111 1111	0000 0000 ... 0000 0000	± Infinity	
0/1	1111 1110	1111 1111 ... 1111 1111	± 1.11..1 x 2 ²⁵⁴⁻¹²⁷	~ ± 3.40 x 10 ³⁸
0/1	1111 1110	1111 1111 ... 1111 1110	± 1.11..0 x 2 ²⁵⁴⁻¹²⁷	
...	
0/1	0000 0010	0000 0000 ... 0000 0000	± 1.00..0 x 2 ²⁻¹²⁷	
0/1	0000 0001	1111 1111 ... 1111 1111	± 1.11..1 x 2 ¹⁻¹²⁷	
...	
0/1	0000 0001	0000 0000 ... 0000 0001	± 1.00..1 x 2 ¹⁻¹²⁷	
0/1	0000 0001	0000 0000 ... 0000 0000	± 1.00..0 x 2 ¹⁻¹²⁷	~ ± 1.18 x 10 ⁻³⁸
0/1	0000 0000	1111 1111 ... 1111 1111	Denormalized number	
0/1	0000 0000	1111 1111 ... 1111 1110		
...		
0/1	0000 0000	0000 0000 ... 0000 0001		
0/1	0000 0000	0000 0000 ... 0000 0000	0	0



Denormalized Numbers

- Assume significand bit = 3 instead of 23 for simplicity:

0 00000010 001 1.001×2^{-125}

0 00000010 000 1.0×2^{-125}

0 00000001 111 1.111×2^{-126}

0 00000001 110 1.11×2^{-126}

...

0 00000001 000 1.0×2^{-126}


0 00000000 111 1.111×2^{-127}

0 00000000 110 1.11×2^{-127}


...

0 00000000 001 1.001×2^{-127}

0 00000000 000 0



Denormalized Numbers



- Assume significand bit = 3:

0 00000010 001 1.001×2^{-125}

0 00000010 000 1.0×2^{-125}

0 00000001 111 1.111×2^{-126}

0 00000001 110 1.11×2^{-126}

...

0 00000001 000 1.0×2^{-126}

0 00000000 111 1.111×2^{-127} $\Rightarrow 0.111 \times 2^{-126}$


0 00000000 110 1.11×2^{-127} $\Rightarrow 0.110 \times 2^{-126}$

...

0 00000000 001 1.001×2^{-127} $\Rightarrow 0.001 \times 2^{-126}$

0 00000000 000 0

} Denormalized
Numbers





Denormal Numbers

- Exponent = 000...0 \Rightarrow hidden bit is 0

$$x = (-1)^s \times (0 + \text{Fraction}) \times 2^{1-\text{Bias}}$$

- Smaller than normal numbers
 - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0


$$x = (-1)^s \times (0 + 0) \times 2^{1-\text{Bias}} = \pm 0.0$$

Two representations
of 0.0!

Range of Numbers Represented by IEEE 754




sign	Exponent	Significand		
0/1	1111 1111	1111 1111 ... 1111 1111	NAN	
0/1	1111 1111	...		
0/1	1111 1111	0000 0000 ... 0000 0001		
0/1	1111 1111	0000 0000 ... 0000 0000	$\pm \text{Infinity}$	
0/1	1111 1110	1111 1111 ... 1111 1111	$\pm 1.11..1 \times 2^{254-127}$	$\sim \pm 3.40 \times 10^{38}$
0/1	1111 1110	1111 1111 ... 1111 1110	$\pm 1.11..0 \times 2^{254-127}$	
...	
0/1	0000 0010	0000 0000 ... 0000 0000	$\pm 1.00..0 \times 2^{2-127}$	
0/1	0000 0001	1111 1111 ... 1111 1111	$\pm 1.11..1 \times 2^{1-127}$	
...	
0/1	0000 0001	0000 0000 ... 0000 0001	$\pm 1.00..1 \times 2^{1-127}$	
0/1	0000 0001	0000 0000 ... 0000 0000	$\pm 1.00..0 \times 2^{1-127}$	$\sim \pm 1.18 \times 10^{-38}$
0/1	0000 0000	1111 1111 ... 1111 1111	$\pm 0.11..1 \times 2^{1-127}$	
0/1	0000 0000	1111 1111 ... 1111 1110	$\pm 0.11..0 \times 2^{1-127}$	
...	
0/1	0000 0000	0000 0000 ... 0000 0001	$\pm 0.00..1 \times 2^{1-127}$	$\sim \pm 1.4 \times 10^{-45}$
0/1	0000 0000	0000 0000 ... 0000 0000	0	0




Floating-Point Addition

- Consider a **4-digit decimal** example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2



Floating-Point Addition


- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} + (-1.110_2) \times 2^{-2}$ ($0.5 + -0.4375$)
- 1. Align binary points
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + (-0.111_2) \times 2^{-1}$
- 2. Add significands
 - $1.000_2 \times 2^{-1} + (-0.111_2) \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625




FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- FP adder usually takes **several cycles**
 - Can be pipelined


FP Adder Hardware





Floating-Point Multiplication

- Consider a **4-digit decimal** example
 - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
 - New exponent = $10 + (-5) = 5$
- 2. Multiply significands
 - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- 3. Normalize result & check for over/underflow
 - 1.0212×10^6
- 4. Round and renormalize if necessary
 - 1.021×10^6
- 5. Determine sign of result from signs of operands
 - $+1.021 \times 10^6$




Floating-Point Multiplication

- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} \times (-1.110_2) \times 2^{-2}$ ($0.5 \times (-0.4375)$)
- 1. Add exponents
 - Unbiased: $(-1) + (-2) = -3$
 - Biased: $((-1) + 127) + ((-2) + 127) = (-3) + 254 - 127 = (-3) + 127$
- 2. Multiply significands
 - $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
 - $1.110_2 \times 2^{-3}$ (no change) with no over/underflow
- 4. Round and renormalize if necessary
 - $1.110_2 \times 2^{-3}$ (no change)
- 5. Determine sign: +ve \times -ve \Rightarrow -ve
 - $-1.110_2 \times 2^{-3} = -0.21875$




FP Arithmetic Hardware

- FP multiplier is of similar complexity to FP adder
 - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
 - Addition, subtraction, multiplication, division, reciprocal, square-root
 - FP \leftrightarrow integer conversion
- Operations usually takes several cycles
 - Can be pipelined




FP Instructions in MIPS

- FP hardware is coprocessor 1
 - Adjunct processor that extends the ISA
- Separate FP registers
 - 32 single-precision: \$f0, \$f1, … \$f31
 - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, …
 - Release 2 of MIPS ISA supports 32×64
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - `lwc1`, `ldc1`, `swc1`, `sdc1`
 - e.g., `ldc1 $f8, 32($sp)`



FP Instructions in MIPS

- Single-precision arithmetic
 - `add.s`, `sub.s`, `mul.s`, `div.s`
 - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
 - `add.d`, `sub.d`, `mul.d`, `div.d`
 - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
 - `c.xx.s`, `c.xx.d` (`xx` is `eq`, `lt`, `le`, ...)
 - Sets or clears FP condition-code bit
 - e.g. `c.lt.s $f3, $f4`
- Branch on FP condition code true or false
 - `bc1t`, `bc1f`
 - e.g., `bc1t TargetLabel`



Example


- C code

...

```
int a, b, c;
```

...

```
c = a + b;
```



```
add $s0, $s1, $s2;
```

...


```
float l, m, n
```

...

```
l = m + n;
```



```
add.s $f0, $f1, $f2;
```




Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- **fahr** in \$f12, result in \$f0, literals(5.0, 9.0, 32.0) in global memory space (\$gp, global pointer)
- Compiled MIPS code:

f2c:	lw	c1	\$f16,	const5(\$gp)	// \$f16 = 5.0
	lw	c2	\$f18,	const9(\$gp)	// \$f18 = 9.0
	div	s	\$f16,	\$f16, \$f18	// \$f16 = 5/9
	lw	c1	\$f18,	const32(\$gp)	// \$f18 = 32
	sub	s	\$f18,	\$f12, \$f18	// \$f18 = fahr-32
	mul	s	\$f0,	\$f16, \$f18	// \$f0 = \$f16 * \$f18
	jr		\$ra		



Example for Guard Bit

- Assume 4-bit significand
- Addition: $1.0000 \times 2^0 + 1.1111 \times 2^{-2}$


$$\begin{array}{r} 1.0000 \times 2^0 \\ + 0.01111 \times 2^0 \\ \hline \end{array}$$

1.01111 rounds to 1.1000 – right! (1.0111, otherwise)
↑
g

- Multiplication:

$$\begin{array}{r} 1.1000 \times 2^0 \\ \times 1.0001 \times 2^{-2} \\ \hline 1.10011000 \times 2^{-2} \end{array}$$

normalize 1.10011000 $\times 2^{-1}$ rounds to 1.1010 $\times 2^{-1}$
↑
g



Rounding Modes

- IEEE Standard has four rounding modes:
 - round to nearest even (default) (반올림)
 - example: base 10
 - $3.56\underline{5}1 \Rightarrow 3.57$ $3.56\underline{4}9 \Rightarrow 3.56$
 - $3.56\underline{5}0 \Rightarrow 3.56$ $3.57\underline{5}0 \Rightarrow 3.58$
 - example: base 2
 - $1.01\underline{1}01 \Rightarrow 1.10$ $1.01\underline{0}11 \Rightarrow 1.01$
 - $1.01\underline{1}00 \Rightarrow 1.10$ $1.00\underline{1}00 \Rightarrow 1.00$
 - $1.00\underline{1}0000001 \Rightarrow 1.01$
 - round towards plus infinity
 - round towards minus infinity
 - round towards 0



Rounding Modes

- “round to nearest” minimizes the mean error introduced by rounding
- Round Bit is calculated to the right of guard bit
- Sticky Bit is used to determine whether there are any 1 bits truncated below the guard and round bits



Round Bit

- Bit to the right of guard bit needed for accurate rounding
- Example: $1.0000 \times 2^0 - 1.0001 \times 2^{-2}$
 - guard and round bits shown


$$\begin{array}{r} 1.0000 \times 2^0 \\ - 0.010001 \times 2^0 \\ \hline 0.101111 \times 2^0 \end{array} \quad \text{Result}$$

↑↑
gr

1.01111×2^{-1} Normalize

1.1000×2^{-1} Round


- Without round bit, result is 1.0111×2^{-1}



Sticky Bit

- ‘Round to even’ problems
 - need to know if the answer is exactly even or not
- Keep “sticky” bit (S):
 - S = 1 if any bits are off to the right, otherwise S = 0
- Example: $1.0000 \times 2^0 + 1.0001 \times 2^{-5}$
 - guard, round, and sticky bits shown

$$\begin{array}{r} 1.0000 \times 2^0 \\ + 0.000010 \times 2^0 \\ \hline 1.000010 \times 2^0 \end{array} \quad \begin{array}{l} \text{(0.000010001} \times 2^0 \text{ to be exact)} \\ \text{Result} \\ \uparrow \text{Round to nearest; without S rounds to } 1.0000 \\ \text{Sticky bit} \end{array}$$




Associativity

- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail
 - $(x + y) + z \neq x + (y + z)$


		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38		-1.50E+38
y	1.50E+38	0.00E+00	
z	1.0	1.0	1.50E+38
		1.00E+00	0.00E+00

- Need to validate parallel programs under varying degrees of parallelism



x86 FP Architecture

- Originally based on 8087 FP coprocessor
 - 8 × 80-bit extended-precision registers
 - Used as a push-down stack
 - Registers indexed from top of stack: ST(0), ST(1), ...
- FP values are 32-bit or 64 in memory
 - Converted on load/store of memory operand
 - Integer operands can also be converted on load/store
- Very difficult to generate and optimize code
 - Result: poor FP performance



Streaming SIMD Extension 2 (SSE2)

- Intel adds 144 new instructions to SSE in 2001
- Adds 4×128 -bit registers
 - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
 - 2×64 -bit double precision
 - 4×32 -bit double precision
 - Instructions operate on them simultaneously
 - Single-Instruction Multiple-Data (SIMD)